

# ECE/CS 5510

## Multiprocessor Programming

### Mutual Exclusion

Chapter 2, The Art of  
Multiprocessor Programming  
(Maurice Herlihy & Nir Shavit)



### Mutual Exclusion



- Today we will try to formalize our understanding of mutual exclusion
- We will also use the opportunity to show you how to argue about and prove various properties in an asynchronous concurrent setting

## Mutual Exclusion



In his 1965 paper E. W. Dijkstra wrote:

"Given in this paper is a solution to a problem which, to the knowledge of the author, has been an open question since at least 1962, irrespective of the solvability. [...] Although the setting of the problem might seem somewhat academic at first, the author trusts that anyone familiar with the logical problems that arise in computer coupling will appreciate the significance of the fact that this problem indeed can be solved."

## Mutual Exclusion



- Formal problem definitions
- Solutions for 2 threads
- Solutions for  $n$  threads
- Fair solutions
- Inherent costs

## Warning

- You will never use these protocols
  - *Get over it*
- You are advised to understand them
  - *The same issues show up everywhere*
  - *Except hidden and more complex*

## Why is Concurrent Programming so Hard?

- Try preparing a seven-course banquet
  - *By yourself*
  - *With one friend*
  - *With twenty-seven friends ...*
- Before we can talk about programs
  - *Need a language*
  - *Describing time and concurrency*

## Time

- “Absolute, true and mathematical time, of itself and from its own nature, flows equably without relation to anything external.” (I. Newton, 1689)
- “Time is, like, Nature’s way of making sure that everything doesn’t happen all at once.” (Anonymous, circa 1968)




## Events

- An event  $a_0$  of thread  $A$  is
  - Instantaneous
  - No simultaneous events (break ties)



# Threads


- A *thread A* is (formally) a sequence  $a_0, a_1, \dots$  of events
  - “Trace” model
  - Notation:  $a_0 \rightarrow a_1$  indicates order



## Example Thread Events

- Assign to shared variable
- Assign to local variable
- Invoke method
- Return from method
- Lots of other things ...

# Threads are State Machines



## States

- Thread State
  - Program counter
  - Local variables
- System state
  - Object fields (shared variables)
  - Union of thread states

## Concurrency

- Thread A



## Concurrency

- Thread A



- Thread B




# Interleavings

- Events of two or more threads
  - Interleaved
  - Not necessarily independent (why?)




# Intervals

- An *interval*  $A_0 = (a_0, a_1)$  is
  - Time between events  $a_0$  and  $a_1$






## Intervals may Overlap




## Intervals may be Disjoint




## Precedence

Interval  $A_0$  precedes interval  $B_0$




## Precedence




- Notation:  $A_0 \rightarrow B_0$
- Formally,
  - End event of  $A_0$  before start event of  $B_0$
  - Also called “happens before” or “precedes”

## Precedence Ordering



- Remark:  $A_0 \rightarrow B_0$  is just like saying
  - 1066 AD  $\rightarrow$  1492 AD,
  - Middle Ages  $\rightarrow$  Renaissance,
- Oh wait,
  - what about this week vs this month?

## Precedence Ordering



- Never true that  $A \rightarrow A$
- If  $A \rightarrow B$  then not true that  $B \rightarrow A$
- If  $A \rightarrow B$  &  $B \rightarrow C$  then  $A \rightarrow C$
- Funny thing:  $A \rightarrow B$  &  $B \rightarrow A$  might both be false!

## Partial Orders

(you may know this already)

- Irreflexive:
  - Never true that  $A \rightarrow A$
- Antisymmetric:
  - If  $A \rightarrow B$  then not true that  $B \rightarrow A$
- Transitive:
  - If  $A \rightarrow B$  &  $B \rightarrow C$  then  $A \rightarrow C$

## Total Orders


(you may know this already)

- Also
  - Irreflexive
  - Antisymmetric
  - Transitive
- Except that for every distinct  $A, B$ ,
  - Either  $A \rightarrow B$  or  $B \rightarrow A$
  - “Totality”

Total order = partial order + totality

## Repeated Events

```
while (mumble) {  
    a0; a1;  
}
```



## Implementing a Counter

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        temp = value;  
        value = temp + 1;  
        return temp;  
    }  
}
```


Make these steps  
*indivisible* using  
locks

## Locks (Mutual Exclusion)

```
public interface Lock {  
    public void lock();  
    public void unlock();  
}
```

## Locks (Mutual Exclusion)

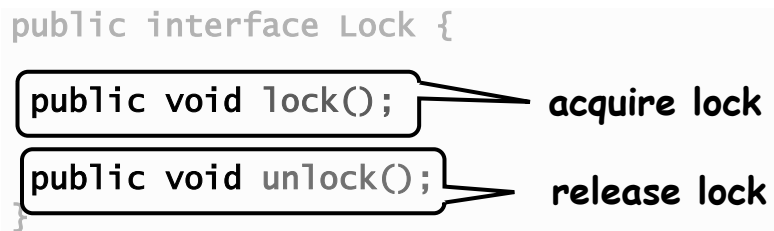
```
public interface Lock {  
    public void lock();  
    public void unlock();  
}
```



A diagram showing a callout box around the `lock()` method signature. An arrow points from the text **acquire lock** to the `lock()` method.

## Locks (Mutual Exclusion)

```
public interface Lock {  
    public void lock();  
    public void unlock();  
}
```



The diagram shows the `Lock` interface with two methods. A callout box points to `lock()` with the text "acquire lock". Another callout box points to `unlock()` with the text "release lock".

## Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```

## Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```

**lock.lock();** → **acquire Lock**

## Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```

**lock.unlock();** → **Release lock  
(no matter what)**



## Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```

Critical section

## Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```

Critical section

*Now, we will formalize properties that a good lock algorithm must exhibit*

## Mutual Exclusion

- Let  $CS_i^k \Leftrightarrow$  be thread  $i$ 's  $k$ -th critical section execution

## Mutual Exclusion


- Let  $CS_i^k \Leftrightarrow$  be thread  $i$ 's  $k$ -th critical section execution
- And  $CS_j^m \Leftrightarrow$  be thread  $j$ 's  $m$ -th critical section execution

## Mutual Exclusion

- Let  $CS_i^k \Leftrightarrow$  be thread  $i$ 's  $k$ -th critical section execution
- And  $CS_j^m \Leftrightarrow$  be  $j$ 's  $m$ -th execution
- Then either
  - $\Leftrightarrow \Leftrightarrow$  or  $\Leftrightarrow \Leftrightarrow$

## Mutual Exclusion

- Let  $CS_i^k \Leftrightarrow$  be thread  $i$ 's  $k$ -th critical section execution
- And  $CS_j^m \Leftrightarrow$  be  $j$ 's  $m$ -th execution
- Then either
  - $\Leftrightarrow \Leftrightarrow$  or  $\Leftrightarrow \Leftrightarrow$



## Mutual Exclusion

- Let  $CS_i^k \leftrightarrow$  be thread  $i$ 's  $k$ -th critical section execution
- And  $CS_j^m \leftrightarrow$  be  $j$ 's  $m$ -th execution
- Then either

-  $\leftrightarrow \leftrightarrow$  or  $\leftrightarrow \leftrightarrow$

$CS_i^k \rightarrow CS_j^m$

$CS_j^m \rightarrow CS_i^k$

## Deadlock-Free



- If some thread calls **lock()**, then some thread will succeed, meaning:
- If some thread calls **lock()**
  - And never returns
  - Then other threads must complete **lock()** and **unlock()** calls infinitely often
- System as a whole makes progress
  - Even if individuals starve

## Starvation-Free



- If some thread calls lock()
  - It will eventually return
- Individual threads make progress
- Note: starvation-freedom implies deadlock-freedom
  - So if former holds, no need to prove latter

## Two-Thread vs $n$ -Thread Solutions

- Two-thread solutions first
  - Illustrate most basic ideas
  - Fits on one slide
- Then  $n$ -Thread solutions

## Two-Thread Conventions

```
class ... implements Lock {  
    ...  
    // thread-local index, 0 or 1  
    public void lock() {  
        int i = ThreadID.get();  
        int j = 1 - i;  
        ...  
    }  
}
```

## Two-Thread Conventions

```
class ... implements Lock {  
    ...  
    // thread-local index, 0 or 1  
    public void lock() {  
        int i = ThreadID.get();  
        int j = 1 - i;  
        ...  
    }  
}
```


Henceforth: *i* is current  
thread, *j* is other thread

## LockOne

```
class LockOne implements Lock {  
    private boolean[] flag =  
                                new boolean[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

## LockOne

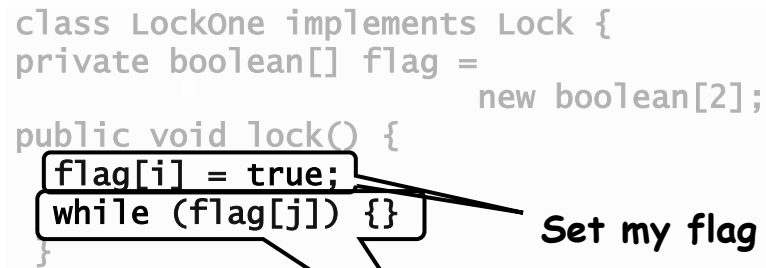
```
class LockOne implements Lock {  
    private boolean[] flag =  
                                new boolean[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```



Set my flag

## LockOne

```
class LockOne implements Lock {  
    private boolean[] flag =  
        new boolean[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

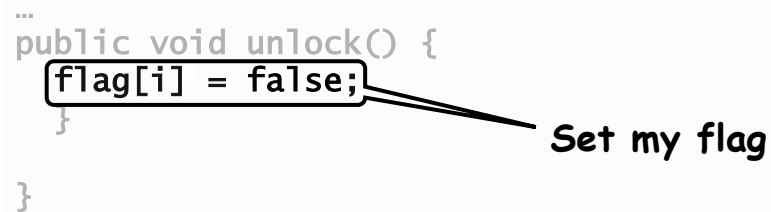


**Set my flag**

**Wait for other  
flag to go false**

## LockOne

```
...  
public void unlock() {  
    flag[i] = false;  
}  
}
```



**Set my flag**



## LockOne Satisfies Mutual Exclusion

- Assume  $CS_A^j$  overlaps  $CS_B^k$
- Consider each thread's last (j-th and k-th) read and write in the lock() method before entering
- Derive a contradiction

## From the Code

- $\text{write}_A(\text{flag}[A]=\text{true}) \rightarrow$   
 $\text{read}_A(\text{flag}[B] \neq \text{false}) \rightarrow CS_A$
- $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow$   
 $\text{read}_B(\text{flag}[A] \neq \text{false}) \rightarrow CS_B$

```
class LockOne implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

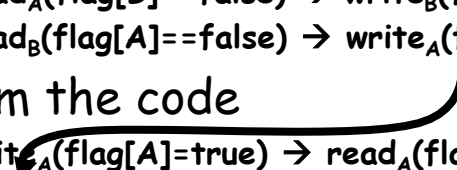
## From the Assumption

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

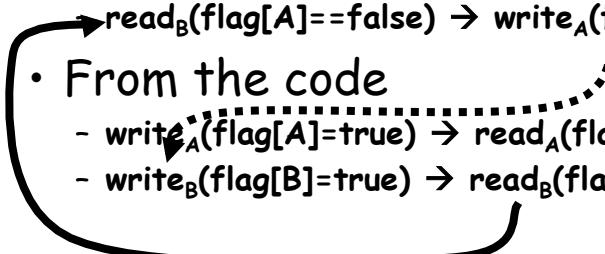
## Combining

- Assumptions:
  - $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
  - $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$
- From the code
  - $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
  - $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

## Combining

- Assumptions:
    - $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
    - $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$
  - From the code
    - $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
    - $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$
- 

## Combining

- Assumptions:
    - $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
    - $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$
  - From the code
    - $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
    - $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$
- 

## Combining

- Assumptions:

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$

- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

- From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$

- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

## Combining

- Assumptions:

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$

- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

- From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$

- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

## Combining


- Assumptions:

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

- From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

Cycle!



# Deadlock Freedom

- LockOne Fails deadlock-freedom
  - Concurrent execution can deadlock

```
flag[i] = true;    flag[j] = true;
while (flag[j]){   while (flag[i]){
```

- Sequential executions OK

## LockTwo

```
public class LockTwo implements Lock {
    private int victim;
    public void lock() {
        victim = i;
        while (victim == i) {};
    }

    public void unlock() {}
}
```

## LockTwo

```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {};  
    }  
  
    public void unlock() {}  
}
```

**Let other go first**

## LockTwo

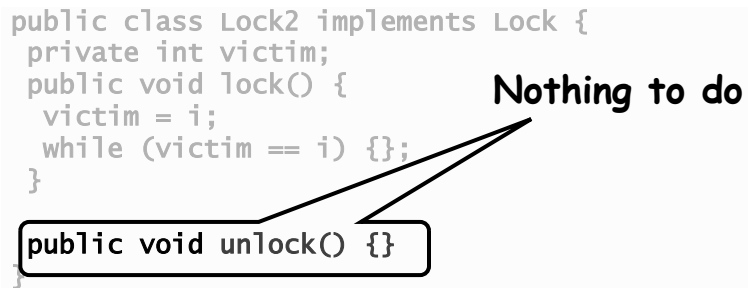
```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {};  
    }  
  
    public void unlock() {}  
}
```

**Wait for permission**

## LockTwo

```
public class Lock2 implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {};  
    }  
    public void unlock() {}  
}
```

**Nothing to do**



## LockTwo Claims

- Satisfies mutual exclusion

- If thread *i* in CS
- Then `victim == i`
- Cannot be both 0 and 1

```
public void LockTwo() {  
    victim = i;  
    while (victim == i) {};  
}
```

- Not deadlock free

- Sequential execution deadlocks
- Concurrent execution does not



## Peterson's Algorithm

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```

## Peterson's Algorithm

**Announce I'm interested**

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```

## Peterson's Algorithm

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```

Announce I'm interested

Defer to other

## Peterson's Algorithm

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```

Announce I'm interested

Defer to other

Wait while other interested & I'm the victim

## Peterson's Algorithm

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```

Announce I'm interested

Defer to other

Wait while other interested & I'm the victim

No longer interested

## Mutual Exclusion

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}
```

- If thread 0 in critical section,
  - flag[0]=true,
  - victim = 1
- If thread 1 in critical section,
  - flag[1]=true,
  - victim = 0

Cannot both be true

## Deadlock Free

```
public void lock() {  
    ...  
    while (flag[j] && victim == i) {};  
}
```

- Thread blocked
  - only at while loop
  - only if it is the victim
- One or the other must not be the victim

## Starvation Free

- Thread *i* blocked  
only if *j* repeatedly  
re-enters so that

flag[j] == true and  
victim == i


- When *j* re-enters
  - it sets victim to *j*.
  - So *i* gets in

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
  
public void unlock() {  
    flag[i] = false;  
}
```

# The Filter Algorithm for $n$ Threads

There are  $n-1$  “waiting rooms” called levels


- At each level
  - At least one enters level
  - At least one blocked if many try
- Only one thread makes it through



## Filter

```
class Filter implements Lock {
    int[] level; // level[i] for thread i
    int[] victim; // victim[L] for level L
```

```
    public Filter(int n) {
        level = new int[n];
        victim = new int[n];
        for (int i = 0; i < n; i++) {
            level[i] = 0;
        }
        ...
    }
```



Thread 2 at level 4

# Filter

```
class Filter implements Lock {
    ...

    public void lock(){
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while (( $\exists$  k != i level[k] >= L) &&
                    victim[L] == i );
        }
    }
    public void unlock() {
        level[i] = 0;
    }
}
```

# Filter

```
class Filter implements Lock {
    ...

    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while (( $\exists$  k != i) level[k] >= L) &&
                    victim[L] == i);
        }
    }
    public void release(int i) {
        level[i] = 0;
    }
}
```

One level at a time

# Filter

```
class Filter implements Lock {  
    ...  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists k \neq i$ ) level[k] >= L) &&  
                victim[L] == i;  
        }  
    }  
    public void release(int i)  
        level[i] = 0;  
}
```

Announce  
intention to  
enter level L

# Filter

```
class Filter implements Lock {  
    int level[n];  
    int victim[n];  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists k \neq i$ ) level[k] >= L) &&  
                victim[L] == i;  
        }  
    }  
    public void release(int i)  
        level[i] = 0;  
}
```

Give priority to  
anyone but me

## Filter

Wait as long as someone else is at same or higher level, and I'm designated victim

```
public void lock() {  
    for (int L = 1; L < n; L++) {  
        level[i] = L;  
        victim[L] = i;  
        while (( $\exists$  k != i) level[k] >= L) &&  
            victim[L] == i);  
    }  
}  
public void release(int i) {  
    level[i] = 0;  
}
```

## Filter


```
class Filter implements Lock {  
    int level[n];  
    int victim[n];  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists$  k != i) level[k] >= L) &&  
                victim[L] == i);  
        }  
    }  
}
```

Thread enters level L when it completes the loop



## Claim

- Start at level  $L=0$
- At most  $n-L$  threads enter level  $L$
- Mutual exclusion at level  $L=n-1$



## Proof is by induction on $L$

- Base case:  $L = 0$
- From code inspection,  $n$  threads are at level  $L = 0$


```

public void lock() {
    for (int L = 1; L < n; L++) {
        level[i] = L;
        victim[L] = i;
        while ((∃ k != i) level[k] == L
            && victim[L] == i) {}
    }
}


public Filter(int n) {
    ...
    for (int i = 0; i < n; i++) {
        level[i] = 0;
    }
    ...
}
    
```

## Induction Hypothesis

- No more than  $n-L+1$  at level  $L-1$
- Induction step: by contradiction
- Assume all at level  $L-1$  enter level  $L$
- A is last to write  $\text{victim}[L]$
- B is any other thread at level  $L$



## Proof Structure



Show that **A** must have seen **B** in  $\text{level}[L]$  and since  $\text{victim}[L] == A$  could not have entered

## From the Code

(1)  $\text{write}_B(\text{level}[B]=L) \rightarrow \text{write}_B(\text{victim}[L]=B)$

```
public void lock() {  
    for (int L = 1; L < n; L++) {  
        level[i] = L;  
        victim[L] = i;  
        while (( $\exists k \neq i$ ) level[k] >= L  
            && victim[L] == i) {};  
    }  
}
```

## From the Code

(2)  $\text{write}_A(\text{victim}[L]=A) \rightarrow \text{read}_A(\text{level}[B])$

```
public void lock() {  
    for (int L = 1; L < n; L++) {  
        level[i] = L;  
        victim[L] = i;  
        while (( $\exists k \neq i$ ) level[k] >= L  
            && victim[L] == i) {};  
    }  
}
```

## By Assumption

(3)  $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$

By assumption, A is the last  
thread to write **victim[L]**

## Combining Observations

(1)  $\text{write}_B(\text{level}[B]=L) \rightarrow \text{write}_B(\text{victim}[L]=B)$

(3)  $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$

(2)  $\text{write}_A(\text{victim}[L]=A) \rightarrow \text{read}_A(\text{level}[B])$

## Combining Observations

(1)  $\text{write}_B(\text{level}[B]=L) \rightarrow$

(3)  $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$

(2)   $\rightarrow \text{read}_A(\text{level}[B])$

```
public void lock() {  
    for (int L = 1; L < n; L++) {  
        level[i] = L;  
        victim[L] = i;  
        while (( $\exists k \neq i$ ) level[k] >= L  
            && victim[L] == i) {};  
    }  
}
```

## Combining Observations

(1)  $\text{write}_B(\text{level}[B]=L) \rightarrow$

(3)  $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$

(2)   $\rightarrow \text{read}_A(\text{level}[B])$

Thus,  $A$  read  $\text{level}[B] \geq L$ ,  
 $A$  was last to write  $\text{victim}[L]$ ,  
so it could not have entered level  $L$ !

## No Starvation

- Filter Lock satisfies properties:
  - Just like Peterson Alg at any level
  - So no one starves
- But what about fairness?
  - Threads can be overtaken by others
  - E.g., thread A which called lock() earlier than B may write to victim[L] later than B

## Bounded Waiting

- Want stronger fairness guarantees
- Thread not “overtaken” too much
- Need to adjust definitions ....

## Bounded Waiting

- Divide `lock()` method into 2 parts:
  - Doorway interval:
    - Written  $D_A$  (for thread A)
    - always finishes in finite steps (no waiting!)
  - Waiting interval:
    - Written  $W_A$  (for thread A)
    - may take unbounded steps

## $r$ -Bounded Waiting

- For threads A and B:
  - If  $D_A^k \rightarrow D_B^j$ 
    - A's k-th doorway precedes B's j-th doorway
  - Then  $CS_A^k \rightarrow CS_B^{j+r}$ 
    - A's k-th critical section precedes B's (j+r)-th critical section
    - B cannot overtake A by more than r times
- First-come-first-served means  $r = 0$ .

*Can we define bounded waiting without doorway/waiting intervals?*

## Fairness Again

- Filter Lock satisfies properties:
  - No one starves
  - But very weak fairness
    - Not  $r$ -bounded for any  $r$ !
  - That's pretty lame...

## Bakery Algorithm

- Provides First-Come-First-Served
- How?
  - Take a "number"
  - Wait until lower numbers have been served
- Lexicographic order
  - $(a,i) > (b,j)$ 
    - If  $a > b$ , or  $a = b$  and  $i > j$



# Bakery Algorithm

```
class Bakery implements Lock {
    boolean[] flag;
    Label[] label;
    public Bakery (int n) {
        flag = new boolean[n];
        label = new Label[n];
        for (int i = 0; i < n; i++) {
            flag[i] = false; label[i] = 0;
        }
    }
    ...
}
```

# Bakery Algorithm

```
class Bakery implements Lock {
    boolean[] flag;
    Label[] label;
    public Bakery (int n) {
        flag = new boolean[n];
        label = new Label[n];
        for (int i = 0; i < n; i++) {
            flag[i] = false; label[i] = 0;
        }
    }
    ...
}
```

Diagram illustrating the Bakery Algorithm state:

- Indices: 0, 2, 6, ..., n-1
- Flag Array: f, f, t, f, f, t, f, f
- Label Array: 0, 0, 4, 0, 0, 5, 0, 0
- Arrows point from the 't' in the flag array at index 2 and the 't' at index 6 to the label array at indices 2 and 6 respectively.
- Below the label array, the letters **CS** are shown.

# Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( (∃ k != i) flag[k]  
                && (label[i],i) > (label[k],k));  
    }  
}
```

# Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while (∃ k flag[k]  
                && (label[i],i) > (label[k],k));  
    }  
}
```

**Doorway**

# Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i],i) > (label[k],k));  
    }  
}
```

**I'm interested**

# Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i],i) > (label[k],k));  
    }  
}
```

**Take increasing  
label (read labels  
in some arbitrary  
order)**

Threads read labels asynchronously, and in arbitrary order: set of labels seen before picking a new one may have never been in memory at the same time

# Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i],i) > (label[k],k));  
    }
```

**Someone is interested**

# Bakery Algorithm

```
class Bakery implements Lock {  
    boolean flag[n];  
    int label[n];  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i],i) > (label[k],k));  
    }
```

**Someone is interested**

Here also, threads read labels asynchronously, and in arbitrary order: set of labels seen before concluding on thread  $k$ 's presence/absence may have never been in memory at the same time

**With lower (label,i) in lexicographic order**

# Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void unlock() {  
        flag[i] = false;  
    }  
}
```

# Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void unlock() {  
        flag[i] = false;  
    }  
}
```

**No longer interested**

**labels are always increasing**

## No Deadlock

- There is always one thread with earliest label (i.e., unique least label/id pair)
- Ties are impossible (why?)

## First-Come-First-Served

- If  $D_A \rightarrow D_B$  then  $A$ 's label is smaller
- And:
  - $\text{write}_A(\text{flag}[A]) \rightarrow$   
 $\text{write}_A(\text{label}[A]) \rightarrow$   
 $\text{read}_B(\text{label}[A]) \rightarrow$   
 $\text{write}_B(\text{label}[B]) \rightarrow$   
 $\text{read}_B(\text{flag}[A])$
- So  $B$  is locked out while  $\text{flag}[A]$  is true

```
class Bakery implements Lock {  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0],  
                        ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i], i) >  
                (label[k], k));  
    }  
}
```

$DF + FCFS = SF$

## Mutual Exclusion

- Suppose A and B in CS together
- Suppose A has a lower (label[A],A)
- When B entered, it must have seen
  - flag[A] is *false*, or
  - (label[A],A) > (label[B],B)

```
class Bakery implements Lock {  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0],  
                        ...,label[n-1])+1;  
        while (∃k flag[k]  
                && (label[i],i) >  
                (label[k],k));  
    }  
}
```

## Mutual Exclusion

- Labels are strictly increasing so
- B must have seen flag[A] == false

## Mutual Exclusion

- Labels are strictly increasing so
- B must have seen  $\text{flag}[A] == \text{false}$
- $\text{Labeling}_B \rightarrow \text{read}_B(\text{flag}[A]) \rightarrow \text{write}_A(\text{flag}[A]) \rightarrow \text{Labeling}_A$

```
class Bakery implements Lock {  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0],  
            ..., label[n-1])+1;  
        while (∃k flag[k]  
            && (label[i],i) > (label[k],k));  
    }  
}
```

## Mutual Exclusion

- Labels are strictly increasing so
- B must have seen  $\text{flag}[A] == \text{false}$
- $\text{Labeling}_B \rightarrow \text{read}_B(\text{flag}[A]) \rightarrow \text{write}_A(\text{flag}[A]) \rightarrow \text{Labeling}_A$
- Which contradicts the assumption that A has a lower  $(\text{label}[A], A)$



## Bakery Y2<sup>32</sup>K Bug

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while (∃k flag[k]  
                && (label[i],i) > (label[k],k));  
    }  
}
```

## Bakery Y2<sup>32</sup>K Bug

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while (∃k flag[k]  
                && (label[i],i) > (label[k],k));  
    }  
}
```

**Mutex breaks if label[i] overflows**

FCFS also doesn't hold

## Does Overflow Actually Matter?

- Yes
  - Y2K
  - 18 January 2038 (Unix `time_t` rollover)
  - 16-bit counters
- No
  - 64-bit counters
- Maybe
  - 32-bit counters

## Timestamps

- Label variable is really a timestamp
- Need ability to
  - Read others' timestamps
  - Compare them
  - Generate a later timestamp
- Can we do this without overflow?

## Timestamps

- Label variable is really a timestamp
- Need ability to
  - Read others' timestamps
  - Compare them
  - Generate a later timestamp
- Can we do this without overflow?
- Can we do this without waiting?

## The Good News

- One can construct a
  - Wait-free (no mutual exclusion)
  - Concurrent
  - Timestamping system
  - That never overflows

## The ~~Bad~~ News

- One can construct a
  - Wait-free (no mutual exclusion)
  - Concurrent


This part is hard

  - Timestamping system
  - That never overflows

## Instead ...


- We construct a Sequential timestamping system
  - Same basic idea
  - But simpler
- Uses mutex to read & write atomically
- No good for building locks
  - But useful anyway

## Precedence Graphs




- Timestamps form directed graph
- Edge  $x$  to  $y$ 
  - Means  $x$  is later timestamp
  - We say  $x$  dominates  $y$

## Unbounded Counter Precedence Graph




- Timestamping = move tokens on graph
- Atomically
  - read others' tokens
  - move mine
- Ignore tie-breaking for now


## Unbounded Counter Precedence Graph



## Unbounded Counter Precedence Graph




## Two-Thread Bounded Precedence Graph




Edges define the ordering:  $0 < 1$ ;  $1 < 2$ ;  $2 < 0$


## Two-Thread Bounded Precedence Graph



## Two-Thread Bounded Precedence Graph




## Two-Thread Bounded Precedence Graph







## Two-Thread Bounded Precedence Graph $T^2$




## Three-Thread Bounded Precedence Graph?




## Three-Thread Bounded Precedence Graph?




## Graph Composition




# Three-Thread Bounded Precedence Graph $T^3$



# Three-Thread Bounded Precedence Graph $T^3$



## Three-Thread Bounded Precedence Graph $T^3$




## In General

$$T^k = T^2 * T^{k-1}$$

K threads need  $3^k$   
nodes

$$\text{label size} = \log_2(3^k) = 2k$$



## Deep Philosophical Question

- The Bakery Algorithm is
  - Succinct,
  - Elegant, and
  - Fair.
- Q: So why isn't it practical?
- A: Well, you have to read N distinct variables

## Shared Memory

- Shared read/write memory locations called Registers (historical reasons)
- Come in different flavors
  - Multi-Reader-Single-Writer (Flag[])
  - Multi-Reader-Multi-Writer (Victim[])
  - Not that interesting: SRMW and SRSW


## Theorem

At least  $N$  MRSW (multi-reader/  
single-writer) registers are needed  
to solve deadlock-free mutual  
exclusion.

$N$  registers like `Flag[]...`


## Proving Algorithmic Impossibility

- To show no algorithm exists:
  - assume by way of contradiction  
one does,
  - in our case assume an alg for  
deadlock free mutual exclusion  
using  $< N$  registers
  - show a bad execution that  
violates properties



## Proof: Need N-MRSW Registers

Each thread must write to some register



...can't tell whether A is in critical section

## Upper Bound

- Bakery algorithm
  - Uses  $2N$  MRSW registers
- So the bound is (pretty) tight
- But what if we use MRMW registers?
  - Like `victim[]` ?

## Bad News Theorem

At least  $N$  MRMW multi-reader/  
multi-writer registers are needed  
to solve deadlock-free mutual  
exclusion.

**(So multiple writers don't help)**


## Theorem (First 2-Threads)

Theorem: Deadlock-free mutual  
exclusion for 2 threads requires at  
least 2 multi-reader multi-writer  
registers

Proof: assume one register suffices  
and derive a contradiction




## Two Thread Execution



- Threads run, reading and writing R
- Deadlock free so at least one gets in

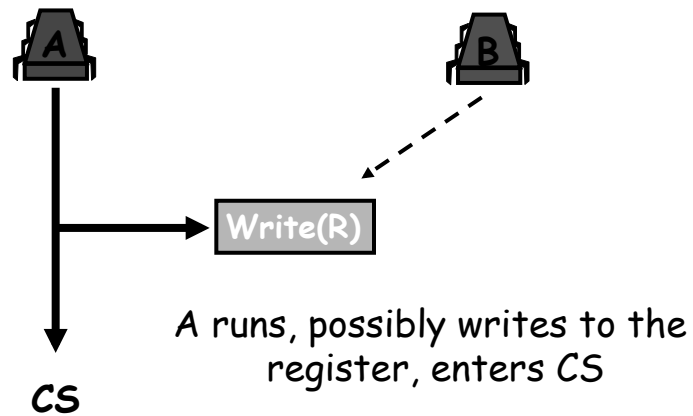
## Covering State for One Register Always Exists

*Before poising to write, B has read register and concluded that CS is empty*

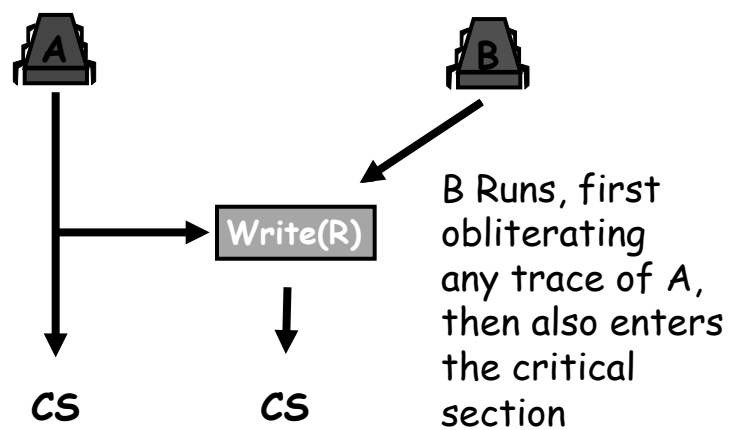


In any mutual exclusion protocol, B has to write to the register before entering CS, so stop it just before

## Proof: Assume Cover of 1




## Proof: Assume Cover of 1




## Theorem

Deadlock-free mutual exclusion for 3 threads requires at least 3 multi-reader multi-writer registers


Proof: Assume Cover of 2



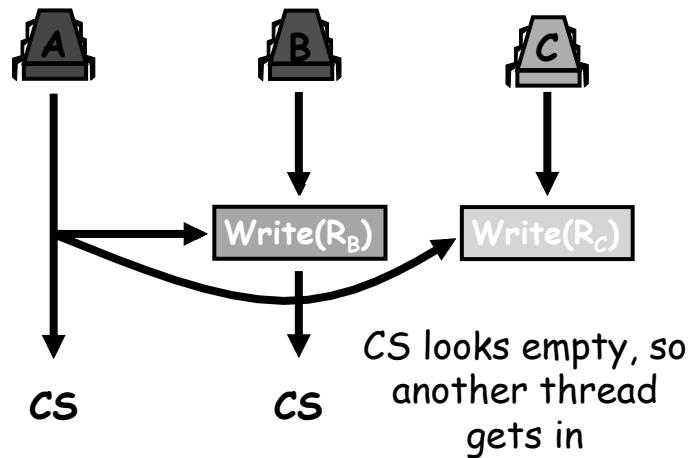
## Run A Solo



## Obliterate Traces of A




## Mutual Exclusion Fails



## Proof Strategy


- Proved: a contradiction starting from a covering state for 2 registers
- Claim: a covering state for 2 registers is reachable from any state where CS is empty

## Covering State for Two




- If we run B through CS 3 times, B must return twice to cover some register, say  $R_B$

## Covering State for Two




- Start with B covering register  $R_B$  for the 1<sup>st</sup> time
- Run A until it is about to write to uncovered  $R_A$ 
  - A must write to  $R_A$  (inconsistent state if it only writes to  $R_B$ )
- Are we done?

## Covering State for Two




- NO! A could have written to  $R_B$
- So CS no longer looks empty

## Covering State for Two



- Run B obliterating traces of A in  $R_B$
- Run B again until it is about to write to  $R_B$
- Now we are done

## Inductively We Can Show



- There is a covering state
  - Where  $k$  threads not in CS cover  $k$  distinct registers
  - Proof follows when  $k = N-1$

## Summary of Lecture

- In the 1960's many incorrect solutions to starvation-free mutual exclusion using RW-registers were published...
- Today we know how to solve FIFO  $N$  thread mutual exclusion using  $2N$  RW-Registers



## Summary of Lecture

- N RW-Registers inefficient
  - Because writes “cover” older writes
- Need stronger hardware operations
  - that do not have the “covering problem”
- In next lectures - understand what these operations are...