

100 points total.

1 Problems (56 points)

1.1 Set (1) [8 points]

In listing 3, two entries are locked before key presence is determined. If no entries were locked and it instead returned `true/false` based on key existence, would this alternative still be linearizable? If so, explain; if not, give a counterexample.

1.2 Set (2) [8 points]

Will listing 1, by itself, function correctly if we switch the locking order on line 9? What about in context with listings 2 and 3? If it has issues in context, how can that be fixed?

1.3 Set (3) [8 points]

Show that listing 1 only needs to lock `pred`.

1.4 Lock-free Linked List (1) [8 points]

Listings 4 and 5 show the `add()` and `remove()` methods for `LockFreeList` class. Describe the two execution scenarios where the `compareAndSet()` operation in line 11 of `add()` method fails because of:

1. expected reference (1st argument)
2. marked value (3rd argument)

1.5 Lock-free Linked List (2) [8 points]

Listings 4 and 5 show the `add()` and `remove()` methods for `LockFreeList` class. Describe the two execution scenarios where the `compareAndSet()` operation in line 11 of `remove()` method fails because of:

1. expected reference (1st argument)
2. marked value (3rd argument)

```
1 public boolean add(T item) {
2     int key = item.hashCode();
3     while (true) {
4         Node pred = head;
5         Node curr = pred.next;
6         while (curr.key < key) {
7             pred = curr; curr = curr.next;
8         }
9         pred.lock(); curr.lock();
10        try {
11            if (validate(pred, curr)) {
12                if (curr.key == key) {
13                    return false;
14                } else {
15                    Node node = new Node(item);
16                    node.next = curr;
17                    pred.next = node;
18                    return true;
19                }
20            }
21        } finally {
22            pred.unlock(); curr.unlock();
23        }
24    }
25 }
```

Listing 1: Add item

```
1 public boolean remove(T item) {
2     int key = item.hashCode();
3     while (true) {
4         Node pred = head;
5         Node curr = pred.next;
6         while (curr.key < key) {
7             pred = curr; curr = curr.next;
8         }
9         pred.lock(); curr.lock();
10        try {
11            if (validate(pred, curr)) {
12                pred.next = curr.next;
13                return true;
14            } else {
15                return false;
16            }
17        }
18    } finally {
19        pred.unlock(); curr.unlock();
20    }
21 }
```

Listing 2: Remove item

```
1 public boolean contains(T item) {
2     int key = item.hashCode();
3     while (true) {
4         Node pred = this.head;
5         Node curr = pred.next;
6         while (curr.key < key) {
7             pred = curr; curr = curr.next;
8         }
9         pred.lock(); curr.lock();
10        try {
11            if (validate(pred, curr)) {
12                return curr.key == key;
13            }
14        } finally {
15            pred.unlock(); curr.unlock();
16        }
17    }
18 }
```

Listing 3: Containment check

```
1 public boolean add(T item) {
2     int key = item.hashCode();
3     while (true) {
4         Window window = find(head, key);
5         Node pred = window.pred, curr = window.curr;
6         if (curr.key == key) {
7             return false;
8         } else {
9             Node node = new Node(item);
10            node.next = new AtomicMarkableReference(curr
11                , false);
12            if (pred.next.compareAndSet(curr, node,
13                false, false)) {
14                return true;
15            }
16        }
17    }
18 }
```

Listing 4: LockFreeList add()

```
1 public boolean remove(T item) {
2     int key = item.hashCode();
3     boolean snip;
4     while (true) {
5         Window window = find(head, key);
6         Node pred = window.pred, curr = window.curr;
7         if (curr.key != key) {
8             return false;
9         } else {
10            Node succ = curr.next.getReference();
11            snip = curr.next.compareAndSet(succ, succ,
12                false, true);
13            if (!snip)
14                continue;
15            pred.next.compareAndSet(curr, succ, false,
16                false);
17            return true;
18        }
19    }
20 }
```


Listing 5: LockFreeList remove()

1.6 Lock-free Linked List (3) [8 points]


A programmer notices that the `add()` method of the `LockFreeList` (Listing 4) never finds a marked node with the same key. Therefore, to save the need to insert a new node, he/she modifies algorithm to simply insert new added object into the existing marked node with same key if such a node exists in the list. Explain why will this not work.

1.7 Compare-And-Swap [8 points]

Some naïve approaches to lock-free deletion in linked lists use a single CAS, swapping the element to be deleted with its next node as shown below.



Unfortunately, this can lead to the following situation:



Why can this happen? Why is it a problem?

2 Programming (44 points)

In this implementation, you will write a microbenchmark to test each of the linked-list algorithms for sets described in chapter 9 of the textbook; for the locking algorithms, replace usage of explicit locks with **synchronized** blocks that provide the same semantics as the explicit lock usage (this will not be possible in all cases). Don't forget to add **volatile** to those variables that need it (but keep usage to a minimum as otherwise performance may suffer). Note that there may be errors in the code that you will need to correct as well.

Use sets of integers `Set<Integer>` as the implementation to benchmark. A worker thread does an operation `ITER` times. The operation on the set could be `add`, `remove` or `contains` with a number passed as a parameter.

ITER = 1000 and operating with numbers (a random choice) from 0 to 100 is a good choice for a worker thread as it ensures some degree of contention.

You will need to vary workload parameters (e.g. ratio of `contains`/`add`/`remove` operations) and find trends in behavior of all algorithms. Vary percentage of `contains` operation from 20% to 80% with the following spacing - 20%, 40%, 60% and 80%. Divide rest of the operations between `add` and `remove` equally wherever possible. Benchmark from 4 threads to 40 threads with reasonable spacing and at least for 8 different thread counts (Use Rlogin).

There are two ways for approaching mixed workloads. Choose whichever you feel comfortable with. (1) Have worker threads that only perform one operation. For example, for 20 threads, you can have 12 threads that always invoke `contains`, and 4 threads each that always invoke `add` and `remove`. Approximate wherever needed. (2) Each thread chooses between `contains`, `add` and `remove` through a uniformly distributed float random number generator.

Hint: Read about `java.util.concurrent.ThreadLocalRandom`. Use it in the `run()` method of your thread.

Don't forget proper benchmarking practices (accounting for JVM warm-up - the more the better, no points of contention beyond the set object itself, etc.); you do not need to use `JMH`, but you may find it useful to avoid the pitfalls discussed in <http://www.oracle.com/technetwork/articles/java/architect-benchmarking-2266277.html>.

Provide the following plots:

1. *throughput vs. threads*; threads varying from 4 to 40 (at least 8 different thread counts); `contains % = 20`; all algorithms.
2. *throughput vs. threads*; threads varying from 4 to 40 (at least 8 different thread counts); `contains % = 40`; all algorithms.
3. *throughput vs. threads*; threads varying from 4 to 40 (at least 8 different thread counts); `contains % = 60`; all algorithms.
4. *throughput vs. threads*; threads varying from 4 to 40 (at least 8 different thread counts); `contains % = 80`; all algorithms.
5. *throughput vs. contains %*; `contains = 20%, 40%, 60%, 80%` for fixed thread count = 20; all algorithms.

Analyze your findings in a writeup, which should not exceed two pages in length.

Your code and writeup (with filename `hw5_<myPID>.pdf`) should be included in a zip file named `hw5_<myPID>.zip` to be uploaded to Canvas by the announced due date; as with homework 1, your code should be part of a `hw5` package and should be runnable from the command line (document the exact command(s) to use).